

# SPARK I/O AND FILE LAYOUT: PARTITIONING BEST PRACTICES & PARQUET INTRODUCTION

# AGENDA

1. **Spark I/O Basics**
2. **Partitioning Essentials**
3. **Performance Impact**
4. **Best Practices: Size & Strategies**
5. **Best Practices: Pruning & Pitfalls**
6. **Best Practices: Summary & Config**
7. **Parquet Introduction**
8. **Parquet Structure: Layers**
9. **Parquet Structure: Diagram**
10. **Key Takeaways**
11. **When to Use & Next Steps**
12. **Conclusion**

# SPARK I/O BASICS

Apache Spark I/O handles efficient data read/write across distributed systems.

## Key Aspects:

- Distributed cluster processing
- Formats: CSV, JSON, Parquet, ORC
- Storage: HDFS, S3, local
- In-memory optimization

## Core Components:

- DataSource API: Schema inference
- Partitioning: Parallelism
- Shuffling: Data redistribution

# SPARK I/O PIPELINE

## Principles:

- Partitioning: Parallel chunks
- Locality: Compute near data
- Fault Tolerance: Lineage
- Caching: Memory storage



## Operations:

- Read/Write: With partitioning
- Streaming: Real-time
- Connectors: Kafka, etc.

# PARTITIONING ESSENTIALS

## What & Why?

Divides data for parallel processing, boosting performance.

## Types:

- Input: File splits (~128MB)
- Shuffle: Transformations
- Output: By columns/keys

## Strategies:

- Range/Hash/File/Custom

## Mechanics:

- Repartition/Coalesce
- Bucketing for joins
- Pruning: Skip irrelevant data

# PERFORMANCE IMPACT

## **Benefits:**

- Parallelism: Full cluster use
- Resource Efficiency: Even load

## **Risks:**

- Too Many: Overhead
- Skew: Imbalance
- Memory: Overload

## **Mitigation:**

- Query alignment
- Monitor skew

# BEST PRACTICES: SIZE & STRATEGIES

**Optimal Size:** 100MB–1GB per partition (balance overhead/memory).

## **Strategies:**

- Filter columns (date/region)
- Avoid high cardinality
- File example: year/month/day


# BEST PRACTICES: PRUNING & PITFALLS

## Pruning:

- Pushdown filters for I/O savings (50-90% reduction)
- Use in WHERE clauses on partition columns

## Pitfalls:

- Too many: Metadata overhead
- Skew: Use salting
- Small files: Coalesce
- Deep nesting: Limits effectiveness





# BEST PRACTICES: SUMMARY & CONFIG

## Core Guidelines:

- Select columns wisely: Frequent filters
- Size appropriately: 100MB–1GB
- Keep shallow: 2-3 levels max
- Query-driven: Match patterns
- Monitor: Skew, files, times

## Key Configurations:

| Parameter          | Purpose    | Rec.      |
|--------------------|------------|-----------|
| maxPartitionBytes  | Input size | 100MB–1GB |
| shuffle.partitions | Shuffle    | 200–2000  |
| adaptive.enabled   | Dynamic    | true      |



# PARQUET INTRODUCTION

## What is Parquet?

Columnar format for Spark/big data analytics.

## Why Use It?:

- Columnar: Read specific columns
- Compression: 50-95% savings
- Pushdown: File-level filters
- Schema Evolution: Parquet supports backward/forward compatibility, allowing addition, deletion, or renaming of columns without rewriting entire datasets, ideal for evolving data pipelines.

## vs. Row Formats:

| Feature | CSV/JSON | Parquet  |
|---------|----------|----------|
| Storage | Large    | Smaller  |
| Speed   | Slow     | Fast     |
| Schema  | Loose    | Strong   |
| Opt.    | Basic    | Advanced |



# PARQUET STRUCTURE: LAYERS

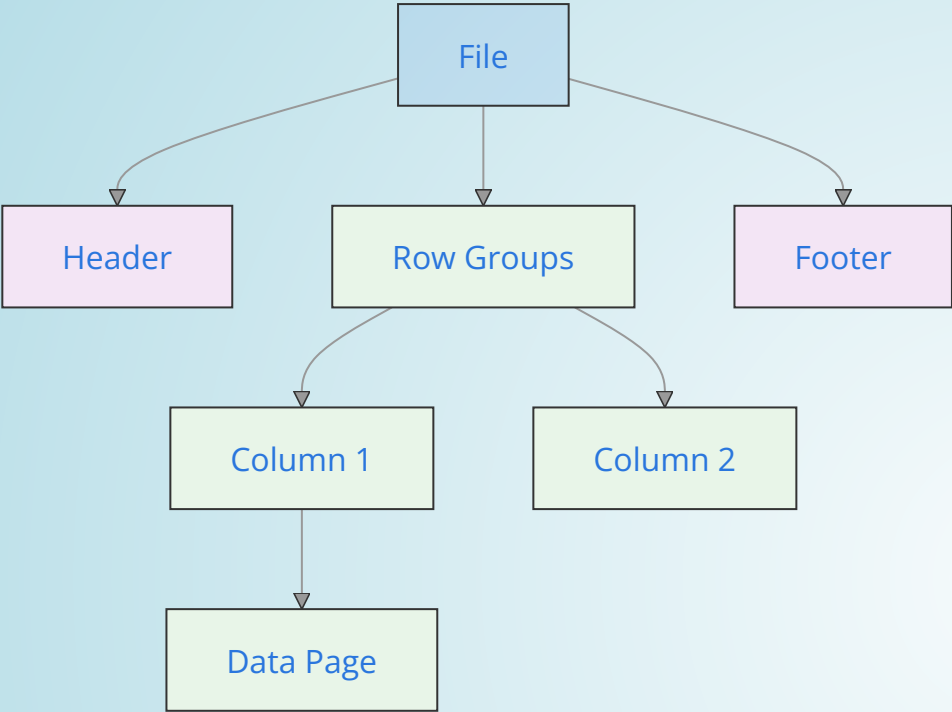
## Key Layers:

1. Header: Magic bytes
2. Row Groups: 128MB–1GB chunks
3. Column Chunks: Per column
4. Pages: ~1MB units
5. Footer: Schema/stats

## Features:

- Encoding: RLE/Dictionary
- Stats: Min/max for pruning
- Types: Nested support

# PARQUET STRUCTURE: DIAGRAM



# KEY TAKEAWAYS

## **Partitioning:**

- Size: 100MB–1GB
- Prune: Filter columns
- Avoid: Skew/high cardinality
- Monitor: Query patterns

## **Parquet:**

- Columnar efficiency
- Compression gains
- 10-100x faster queries
- For analytics/large data

## **Checklist:**

- Analyze patterns
- Optimize sizes/compression
- Test pruning
- Monitor performance

# WHEN TO USE & NEXT STEPS

## Use Cases:

- Analytics/reporting
- Data lakes
- ETL with aggregations
- Evolving Schemas: Ideal for datasets where structure changes over time, as Parquet handles schema modifications efficiently without data loss or full reprocessing.

## Implementation:

- Low-cardinality partitions
- Snappy compression
- Adaptive Spark enabled
- Iterate on metrics

## CONCLUSION

- **Partitioning:** Aim for 100MB–1GB partitions, use low-cardinality columns for pruning, monitor for skew.
- **Parquet:** Columnar storage for efficient analytics, compression, and schema evolution.

**Apply these practices to optimize your Spark workflows and achieve better performance!**

